

Algorithms and Data Structures III (course 1DL481)

Uppsala University – Spring 2023

Report for Assignment 1 by Team 7

Anderson LEONG

Louis HENKEL

12th February 2023

Problem 1: Mixed Integer Programming (MIP)

Task a: Model. We define the variables first:

- Let z be the number of zones.
- Let s be the number of service stations that we place in one zone each.
- Let v be the number of vehicle, we will therefore have $N = v \cdot s$ of vehicles at our disposal in total.
- Let c be the number of vehicles considered
- Let d be a tuple of length z giving the demand for each zone z .
- Let T be a matrix of dimension $z \times z$ with the travel time from one zone to the other.

We create a cost matrix C , which is generated from multiplying each line of the time matrix T_i with the i -th value of the demand tuple. In addition we have a binary vector \mathbf{s} of length z where the i -th value is 1 if we have a station zone i . Finally we have a matrix A of dimension $z \times z$. We have these three constraints:

$$A_{i,j} \in [0, v] \quad \forall i, j \in [1, z] \quad (1)$$

$$\sum_{i=1}^z A_{i,j} = c \quad \forall j \in [1, z] \quad (2)$$

$$\|\mathbf{s}\| = s \quad (3)$$

Our cost function is given by:

$$\text{cost} = \sum_{i=1}^z \sum_{j=1}^z j = 1^z A_{i,j} \cdot C_{i,j}$$

Task b: Implementation. Our model `servStatLoc.mod` is [uploaded](#) with this report: we [checked](#) that its constraints and objective function are linear (and we are aware that four points will otherwise be deducted from our score for this problem). We chose the MIP solver [Gurobi](#) for our experiments, which we ran on the NEOS server for MINTO/AMPL with a CPU time set to 4 hours.

Task c: 10 Zones. The results are in Table 1. When s increases, the optimal objective value decrease, because we have more stations, therefore the average distance is lower, because the new stations may be closer to some of the zones then the initial stations.

Task d: 20 Zones. The results are in Table 1. When s grows beyond 4, the optimal objective value is 0.011920 for $s = 5$.

Task e: 40 Zones. The results are in Table 1.

Task f: 80 Zones. The results are in Table 1. Upon $s = 16$ service stations with $v = 1$ vehicle each, the optimal objective value is 0.013908 the one for $s = 8$ and $v = 2$, because 0.116972.

Task g: 120 Zones. The results are in Table 1. Our model does not time out.

Task h: 250 Zones. The results are in Table 1. Our model does not time out.

Task i: Brute-Force Algorithm. The size of the search space of a totally brute-force search algorithm is

$$\binom{z}{s} = \frac{z!}{s! \cdot (z-s)!} = \frac{z \cdot (z-1) \cdots (z-s+1)}{1 \cdot 2 \cdots s},$$

because if we want to position s objects in z places. The first object we position has z places it can go to, the second one has $z-1$ possibilities, and so on until the last one that $z-s+1$ options. As the order in which we make position the station does not matter, we have to divide by the possibilities in which we can position the station. High school basic combinatorics therefore give us the result above.

The numbers of candidate solutions this brute-force search algorithm has to examine per second in order to match the reported runtime performance of [Gurobi](#) on our model are given in the right-most column of Table 1, for each instance that [Gurobi](#) solved to proven optimality without timing out.

z	s	v	c	time	objective value	optimality gap	brute-force
10	2	2	3	0.00	0.008740	0.00%	∞
10	3	2	3	0.01	0.007145	0.00%	$1, 2 \cdot 10^4$
10	4	2	3	0.01	0.005884	0.00%	$2, 1 \cdot 10^4$
20	2	2	3	0.02	0.023246	0.00%	$9.5 \cdot 10^3$
20	3	2	3	0.02	0.016681	0.00%	$5.7 \cdot 10^4$
20	4	2	3	0.02	0.013908	0.00%	$2.423 \cdot 10^5$
20	5	2	3	0.03	0.011920	0.00%	$5.168 \cdot 10^5$
20	6	2	3	0.01	0.010839	0.00%	$3.876 \cdot 10^6$
40	5	2	3	0.34	0.048470	0.00%	$1.935 \cdot 10^6$
80	8	2	3	8.06	0.124612	0.00%	$3.596 \cdot 10^9$
80	16	1	3	0.06	0.116972	0.00%	$4.493 \cdot 10^{17}$
120	10	2	3	35.94	0.231958	0.00%	$4.848 \cdot 10^{12}$
250	12	3	4	383.36	0.523060	0.00%	$2.482 \cdot 10^{17}$

Table 1: Service station location: runtime (in seconds), objective value, and optimality gap (in percent; positive if an optimal solution was not found and proven before timing out) using [Gurobi](#), with a timeout of 14400 CPU seconds. The right-most column gives the number of candidate solutions the brute-force search algorithm has to examine per second in order to match the runtime performance of [Gurobi](#), if the instance was solved to proven optimality, and ‘n/a’ for ‘non-applicable’ otherwise.

Problem 2: Stochastic Local Search (SLS)

A lower bound on the number λ of shared elements of any pair among v subsets of size r drawn from a given set of b elements is given in [?]:

$$\text{lb}(\lambda) = \left\lceil \frac{\left\lceil \frac{rv}{b} \right\rceil^2 ((rv) \bmod b) + \left\lfloor \frac{rv}{b} \right\rfloor^2 (b - ((rv) \bmod b)) - rv}{v(v-1)} \right\rceil \quad (4)$$

Task a: SLS Algorithm.

1. Representation. The current state of the local search is given by a $v \times b$ matrix, where each row v is a binary array of length b where there are always r elements that are 1.
2. Initial Assignment. It is made by randomly generating v binary arrays where r ones are randomly positioned.
3. Move. Each move we make is a switch in one row, making a permutation of a zero element with a one element.
4. Constraints. Each row must have exactly r elements which are one, which is a built in constrain which is never violated, because we only make swaps in the moves and each row gets initialised this way.
5. Neighbourhood. As for each of the v rows, we create one possible move, there is a total of v moves in the neighbourhood. The search space is connected, as each permutation of r from b elements can be found from a finite sequence of permutations and switches. As we can reach every feasible configuration for 1 row, we can also reach every feasible configuration of the matrix, but there will most likely be moves that increase cost between any given configuration and an optimal solution.
6. Cost Function. Let v_i and v_j be the i -th and j -th row of a configuration. Then our cost function is given by:

$$\text{cost} = \max_{i \neq j} v_i \cdot v_j$$

7. Probing. We iterate through the rows, look at the r elements that have a value 1 and determine, which one is the most expensive one. We look at the $r - v$ elements that have the value 0 and see, which one is the cheapest one. The maximum change in the cost is 1. If we find an swap where we reduce the cost by 1 we make the move. If the cost difference is 0 or -1 , we do not make the move. As the rows are represented in arrays, we can evaluate the cost difference from one element of the neighbourhood, we only need to evaluate the dot product of the new row and all the other rows, so we have a complexity of $O(v \cdot b)$
8. Heuristic. We make sure, that when in a given state, we do not probe the same row twice, so that the same neighbour is not probed twice. The entire neighbourhood does is not explored exhaustively and if found the first move that decreases the cost is selected. Once we have probed each row, we know, that the neighbourhood is exhausted, and can make a random move instead, and add the current state in our tabu list. We decided to do it this way, because if we calculate the cost of each neighbour, it would drastically increase the run time.

9. **Optimality.** As the Equation 4 gives a lower bound for the maximal cost (except of the input of $\langle 10, 8, 3 \rangle$, we have found an optimal solution, if we reached the lower bound given by Equation 4, and can stop the algorithm. If the number of restarts set in the parameter β is reached, we also stop and return the best found configuration up until this point.
10. **Meta-Heuristic.** All moves made before arriving at the suboptimal local minima, will be committed and added to the tabu list for alpha iterations. This prevents the same moves from being made to escape from the local minima and hopefully enter a different neighbourhood.
11. **Random Restarts.** We can make random restarts when we the probing does not yield a better result. The number of restarts is given by the parameter β .

Task b: Implementation. We chose the high-performance programming language C++, for which a compiler is available on the Linux computers of the IT department. All source code is [uploaded](#) with this report. The compilation and running instructions are to have all the .cpp files, .h files, .cc file and the makefile in a folder with the terminal in the folder, enter the commands:

make

and run:

./InvDes v b r

to run the executable generated on the input parameters of choice v , b , r

./InvDes

to run the executable on the pre-input based on the java skeleton code provided.

The executable writes to the standard output a line followed by one line per row of the time of a run as well as the cost. If the argument p is written before the parameter v , b and r , then a space separated matrix representing the solution is also printed.

Task c: Experiments. All experiments were run under the the *gullviva.it.uu.se* server from the IT department, with Ubuntu 22.04 x86_64 architecture.

We attempted to have α be the number of iterations a move will be in the taboo list for, and β be the number of restarts from the original starting board to allow the search to escape the local minima, however, changing the parameters led to increased run time but with little to no improvement in performance. We suspect certain helper functions to be bugged that is causing our taboo search to not function as intended and hence its ability to arrive at the optimal solutions is low, where they only remain in the local minima regardless of the alpha beta parameters, hence, we chose to reduce these values to reduce run time.

Task d: Exact Algorithm. An exact algorithm could work as follows: It could perform brute force search, trying out all possible non-equivalent configurations. Two configurations are equivalent, if a permutations of the rows leads to both configurations to be the same. This would lower the size of the search space. For one row, the search space is given by $\binom{b}{r}$. As we have a total of b rows, there are in total $\binom{b}{r}^v$ possibilities. As the permutation of v elements yield a group of size $v!$, the number of equivalent classes is $b!$. Hence, the total search space is of size

$$\frac{\binom{b}{r}^v}{v!}$$

v	b	r	$\text{lb}(\lambda)$	$\langle \alpha, \beta \rangle = \langle 2, 5 \rangle$			$\langle \alpha, \beta \rangle = \langle 10, 20 \rangle$			exact
				time	steps	λ	time	steps	λ	
10	30	9	2	1.5285438	5	3	5.7164	5	3	$6.4787375 \cdot 10^{64}$
12	44	11	2	6.976366	5	3	19.19752	5	3	$1.239164 \cdot 10^{109}$
15	21	7	2	1.653098	5	3	5.189526	5	3	$4.4440315 \cdot 10^{63}$
16	16	6	2	0.9467204	5	3	62.2835612	5	3	$1.4439218 \cdot 10^{49}$
9	36	12	3	2.148164	5	4	64.446644	5	4	$9.6739127 \cdot 10^{75}$
11	22	10	4	61.01132	5	5	3.00055	5	5	$3.3944265 \cdot 10^{54}$
19	19	9	4	62.539938	5	5	125.586988	5	5	$2.9144089 \cdot 10^{75}$
10	37	14	2	3.79418	5	5	70.73318	5	5	$5.2415527 \cdot 10^{90}$
8	28	14	6	0.761037	5	7	2.896216	5	7	$2.1860830 \cdot 10^{56}$
10	100	30	6	46.0923	5	9	167.7912	5	9	$2.857551 \cdot 10^{246}$
6	50	25	10	2.465378	5	11	122.977204	5	11	$2.2987147 \cdot 10^{81}$
6	60	30	12	4.473836	5	13	68.85049	5	14	$8.4940661 \cdot 10^{98}$
11	150	50	7	257.1908	5	17	301.199	5	17	∞
9	70	35	11	17.96508	5	17	147.90506	5	17	$4.317829 \cdot 10^{173}$
10	350	100	4	306.8834	5	30	308.7114	5	30	∞
13	250	80	18	303.6696	5	27	302.5026	5	28	∞
10	325	100	21	303.907	5	33	305.1178	5	32	∞
15	350	100	19	309.9542	5	37	310.6142	5	36	∞
9	300	100	25	306.5702	5	37	303.6008	5	35	∞
12	200	75	17	301.304	5	30	302.142	5	28	∞
10	360	120	22	306.7936	5	46	307.01	5	46	∞

Table 2: Investment design: median runtime (in seconds), median number of steps, and median achieved λ , for two configurations of values for the local-search parameters α and β , over 5 independent runs per instance, with a timeout of 300.0 CPU seconds per run. The right-most column gives the number of candidate solutions the outlined exact algorithm has to examine per second in order to match the runtime performance of the seemingly best configuration of values for the local-search parameters, namely $\langle \alpha, \beta \rangle = \langle 2, 5 \rangle$, if the instance was solved to proven optimality, and ‘n/a’ for ‘non-applicable’ otherwise.

The number of candidate solutions this exact algorithm has to examine per second in order to match the runtime performance of the seemingly best configuration of values for the local-search parameters, according to Task c, of our stochastic local search algorithm is given in the right-most column of Table 2, for each instance solved to proven optimality.

Feedback to the Teachers

The help sessions were very helpful, as sometimes the task were not very clear. The demo report with the clear structure was very helpful in what to do and what to write. The MIP task involved more time in the development of the model, but once you had one, the rest of the tasks were pretty easy, while the SLS task involved much more time writing the code and coming up with solutions that came along the way. We as a team mostly failed at time management, which is something we would like to improve for the second assignment.